

# COMPLETED: Task 07, dplyr

*Patrick McKenzie*

*28 February 2017*

## Resources

- R Studio's Data Wrangling Cheat Sheet. List of all R Studio cheat sheets here
- R for Data Science: Chapter 5, 9-12
- Wickham (2014) Tidy Data
- The dplyr vignette
- Regular expressions guide - there are many quickstart guides and cheat sheets on the web. I think this one is pretty good.

## Notes on built-in datasets

It makes sense to practice **dplyr** using large-ish data sets, since **dplyr** is designed handle big(ish) data. Because those data sets make the package file size considerably bigger, they are distributed in packages of related data sets. To use the data sets, simply install and load them as you would with a regular package. So, for instance:

```
install.packages("nycflights13")
library(nycflights13)
```

We will be using two data packages: **nycflights13** and **babynames**. **nycflights13** contains the five distinct data sets:


- airlines
- airports
- flights
- planes
- weather

## Tasks

All tasks should be done, to the extent possible, using tidyverse syntax including piping and functions.

- Using the `nycflights13::weather`:
  - Determine whether there are any clear outliers in wind speed (`wind_speed`) that should be rejected. If so, filter those bad point(s) and proceed.
  - What direction has the highest median speed at each airport? Make a table and a plot of median wind speed by direction, for each airport. *Optional fun challenge: If you like, this is a rare opportunity to make use of `coord_polar()`.*

```
plot(nycflights13::weather$wind_speed) #Shows one super-high value
```



```
sum(is.na(nycflights13::weather$wind_speed)) #Shows that we have three NA wind speed values
```

```
## [1] 3
```

```
filtered.weather <- nycflights13::weather %>%
  filter(!(wind_speed > 50 | is.na(wind_speed))) #This removes our outlier and NA values
summarised_wind <- filtered.weather %>%
  group_by(origin,wind_dir) %>%
  summarise(median_wind_speed = median(wind_speed)) #This makes a data frame of median wind speeds by a
head(summarised_wind)
```

```
## Source: local data frame [6 x 3]
## Groups: origin [1]
##
##   origin wind_dir median_wind_speed
##   <chr>    <dbl>         <dbl>
## 1   EWR      0           0.00000
## 2   EWR     10           9.20624
## 3   EWR     20           9.20624
## 4   EWR     30           9.20624
## 5   EWR     40          10.35702
## 6   EWR     50           8.05546
```

```
highest_med_speeds <- summarised_wind %>%
```


```
  group_by(origin) %>%
  filter(median_wind_speed == max(median_wind_speed)) #This selects the direction for each airport at w
highest_med_speeds #Here are the directions with the highest median wind speeds at each airport
```

```
## Source: local data frame [10 x 3]
## Groups: origin [3]
##
##   origin wind_dir median_wind_speed
##   <chr>    <dbl>         <dbl>
## 1   EWR     290          12.65858
```

```
## 2    EWR      300      12.65858
## 3    EWR      320      12.65858
## 4    EWR      330      12.65858
## 5    JFK      290      14.96014
## 6    JFK      300      14.96014
## 7    JFK      310      14.96014
## 8    JFK      330      14.96014
## 9    LGA      270      13.80936
## 10   LGA      290      13.80936
```

```
summarised_wind %>%
  ggplot(aes(x = wind_dir, y = median_wind_speed,color = origin)) + geom_point()
```

```
## Warning: Removed 3 rows containing missing values (geom_point).
```



- Using `nycflights13::flights` and `nycflights13::airlines`:
  - Make a table with two columns: airline name (not carrier code) and median distance flown from JFK airport. The table should be arranged in order of decreasing mean flight distance. Hint: use a `_join` function to join flights and airlines.

```
table <- (nycflights13::flights[,c("carrier","distance")] %>%
  left_join(nycflights13::airlines,by = "carrier"))[, -1] #Makes a tbl of airlines and distances
head(table)
```

```
## # A tibble: 6 × 2
##   distance      name
##   <dbl>      <chr>
## 1    1400 United Air Lines Inc.
## 2    1416 United Air Lines Inc.
```

```
## 3      1089 American Airlines Inc.
## 4      1576      JetBlue Airways
## 5       762    Delta Air Lines Inc.
## 6       719    United Air Lines Inc.
```

```
summ_table <- table %>%
  group_by(name) %>%
  summarise(median.distance = median(distance), mean.distance = mean(distance)) %>%
  arrange(desc(mean.distance))
summ_table
```

```
## # A tibble: 16 × 3
##           name median.distance mean.distance
##           <chr>          <dbl>         <dbl>
## 1    Hawaiian Airlines Inc.      4983      4983.0000
## 2      Virgin America          2475      2499.4822
## 3    Alaska Airlines Inc.      2402      2402.0000
## 4    Frontier Airlines Inc.      1620      1620.0000
## 5    United Air Lines Inc.       1400      1529.1149
## 6    American Airlines Inc.      1096      1340.2360
## 7      Delta Air Lines Inc.      1020      1236.9012
## 8      JetBlue Airways          1023      1068.6215
## 9    Southwest Airlines Co.       748       996.2691
## 10 AirTran Airways Corporation    762       664.8294
## 11      Envoy Air                502       569.5327
## 12 ExpressJet Airlines Inc.       533       562.9917
## 13      US Airways Inc.          529       553.4563
## 14    Endeavor Air Inc.          509       530.2358
## 15    SkyWest Airlines Inc.       419       500.8125
## 16    Mesa Airlines Inc.         229       375.0333
```

- Make a *wide-format* data frame that displays the number of flights that leave Newark (“EWR”) airport each month, from each airline

```
numflights_table <- nycflights13::flights %>%
  left_join(nycflights13::airlines, by = "carrier") %>%
  filter(origin == "EWR") %>%
  group_by(month, name) %>%
  summarise(number.flights = length(flight)) %>%
  spread(month, number.flights)
numflights_table
```


```
## # A tibble: 12 × 13
##           name `1` `2` `3` `4` `5` `6` `7`
## *      <chr> <int> <int> <int> <int> <int> <int> <int>
## 1    Alaska Airlines Inc.    62    56    62    60    62    60    62
## 2    American Airlines Inc.   298   268   295   288   297   291   303
## 3      Delta Air Lines Inc.   279   249   319   364   377   347   340
## 4    Endeavor Air Inc.       82    75    91    88   103    88    94
## 5      Envoy Air            212   196   228   220   226   218   228
## 6 ExpressJet Airlines Inc.  3838  3480  3996  3870  4039  3661  3747
## 7      JetBlue Airways       573   532   612   567   517   506   546
## 8    SkyWest Airlines Inc.    NA    NA    NA    NA    NA     2    NA
## 9    Southwest Airlines Co.   529   490   532   518   530   501   526
## 10   United Air Lines Inc.  3657  3433  3913  4025  3874  3931  4046
## 11      US Airways Inc.      363   328   372   361   381   390   402
```

```
## 12      Virgin America    NA    NA    NA    170    186    180    181
## # ... with 5 more variables: `8` <int>, `9` <int>, `10` <int>, `11` <int>,
## #   `12` <int>
```

- Using the **babynames** dataset:
  - Identify the ten most common male and female names in 2014. Make a plot of their frequency (prop) since 1880. (This may require two separate piped statements).
  - Make a single table of the 26th through 29th most common girls names in the year 1896, 1942, and 2016

```
common_names <- babynames[babynames$year == 2014,] %>% #Selects rows from 2014
  group_by(sex) %>% #Groups by sex so that we get the top group from each in the next line
  top_n(10, n) #This gives us a data frame of top 10 names from 2014.
```

```
babynames %>%
  filter(paste0(sex,name) %in% paste0(common_names$sex,common_names$name)) %>% #selects names for each
  ggplot(aes(x = year, y = prop, color = name)) + geom_line() #Plots frequencies since 1880.
```



```
#####
filtered_df <- babynames %>%
  filter(year %in% c("1896", "1942", "2014")) %>%
  filter(sex == "F") #This pulls out girls' names from the three years

girlsnames <- integer(0)
for (i in c("1896", "1942", "2014")) {
  girlsnames <- rbind(girlsnames, (filtered_df %>%
    filter(year == i))[26:29,]) #this selects rows 26-29 from each year
```

```
}
girlsnames
```

```
## # A tibble: 12 × 5
##   year  sex  name    n    prop
## *   <dbl> <chr>   <chr> <int>   <dbl>
## 1  1896    F  Martha 2022 0.008023969
## 2  1896    F   Esther 1964 0.007793805
## 3  1896    F Frances 1964 0.007793805
## 4  1896    F   Edith 1932 0.007666819
## 5  1942    F   Helen 10014 0.007202575
## 6  1942    F  Marilyn 9904 0.007123458
## 7  1942    F   Diane 9550 0.006868843
## 8  1942    F   Martha 9513 0.006842231
## 9  2014    F Brooklyn 6767 0.003490782
## 10 2014    F    Lily 6727 0.003470148
## 11 2014    F  Hannah 6512 0.003359240
## 12 2014    F   Layla 6428 0.003315908
```


- Write task that involves some of the functions on the Data Wrangling Cheat Sheet and execute it. You may either use your own data or data packages (e.g., the ones listed here).

**Plot average temperature per year for the area within five degrees of latitude/longitude of Knoxville.**

```
atmos_data <- nasaweather::atmos
Kville.lon.lat <- ggmap::geocode("Knoxville")
```

```
## Information from URL : http://maps.googleapis.com/maps/api/geocode/json?address=Knoxville&sensor=fal
```

```
Kville_area_data <- atmos_data %>%
  filter(long > (Kville.lon.lat$lon-5) & long < (Kville.lon.lat$lon+5) &
         lat > (Kville.lon.lat$lat-5) & lat < (Kville.lon.lat$lat+5))
Kville_area_data %>%
  group_by(year,month) %>%
  summarise(temp = mean(temp)) %>%
  ggplot(aes(x=month,y=temp,color=as.factor(year))) + geom_line(size = 1.5) + scale_colour_manual(values=
```



## Optional challenge

Using regular expressions, make a plot of the change in frequency of some letter pattern in names. For instance: how has the frequency of female names ending in “leigh” changed over time relative to names ending in “lee”?

```
substrRight <- function(x, n){
  substr(x, nchar(x)-n+1, nchar(x))
} #Stole this nice function from StackOverflow http://stackoverflow.com/questions/7963898/extracting-th

leighnames <- babynames %>%
  filter(sex == "F") %>%
  filter(substrRight(name,5) == "leigh") %>%
  group_by(year) %>%
  summarise(leigh = length(name)) %>%
  gather(ending,number,-year) #This builds a tibble of names with "leigh" endings.
babynames %>%
  filter(sex == "F") %>%
  filter(substrRight(name,3) == "lee") %>%
  group_by(year) %>%
  summarise(lee = length(name)) %>%
  gather(ending,number,-year) %>%
  bind_rows(leighnames) %>%
  ggplot(aes(x = year,y = number, color = ending)) + geom_line()
```

